


The Mythos Paradox

Identity-Centric Resilience in the Agentic Era

How Claude Mythos 2 rewrites the threat landscape and why Ping Identity is the only platform built to operate at AI speed.

CONFIDENTIAL | For Security Leadership



EXECUTIVE SUMMARY

The security equation changed in 2026.
Identity is the only durable answer.



The Threat Is Real and It Is Now

Claude Mythos 2 Preview develops working exploits in hours — not weeks. AI-speed attacks have outpaced human-speed patch cycles permanently.



The Attack Surface Has Shifted

As code hardens, adversaries pivot to identity: deepfake impersonation, synthetic fraud, MCP token hijacking, and agent credential abuse.



Ping Identity Is Built for This Moment

Identity for AI and Verified Trust give enterprises the only runtime-enforced, AI-speed identity fabric commercially available today.

THE MYTHOS THREAT IN NUMBERS

Claude Mythos 2 Preview — red-team benchmark data, Q1 2026

181

Working Exploits

vs. ~0 for prior generation

83.1%

CyberGym Score

+16.5 points over Opus 4.6

Hours

Exploit Timeline

Reduced from weeks or months

6–18 mo

Proliferation Window

Before adversaries acquire Mythos-class AI



The Patch Gap Is Permanent

When weaponization compresses from weeks to hours, patch cycles cannot close the window. Mythos chains low-severity CVEs — bypassing KASLR and stack protections — to full system compromise before an analyst finishes triage.



A 6 to 18 Month Strategic Clock

Project Glasswing gives defenders a temporary asymmetric advantage. Analysts estimate adversaries will acquire equivalent Mythos-class capability within 6 to 18 months. The window to build identity-centric resilience is now — not after the first AI-speed breach.

THE IDENTITY PIVOT: WHERE AI ATTACKS LAND

As code hardens, adversaries attack the human layer — and the agentic layer.



EXTERNAL IDENTITY THREATS



AI-Scale Social Engineering

Personalized phishing at industrial scale. Every employee is a high-value target.



Real-Time Deepfake Fraud

Voice and face-swap in live video calls defeats KYC and authorizes fraudulent transactions.



Synthetic Identity Onboarding

AI-generated identities defeat document-only verification at the customer enrollment stage.



INTERNAL AGENTIC THREATS



Agent Credential Impersonation

Agents using human credentials create accountability gaps invisible to SIEM and DLP tools.



MCP Token Hijacking

A single compromised MCP token grants lateral access across the entire agentic pipeline.



Privilege Scope Creep

Agents accumulate permissions autonomously — becoming the equivalent of a compromised domain admin.

THE PING IDENTITY RESPONSE

Four capability domains that operate at AI speed



Agent IAM Core

Registers AI agents as first-class identities with unique credentials and scoped permissions. OAuth 2.0 RFC 8693 Token Exchange issues delegation tokens — agents never hold human credentials. Every action is auditable.



Agent Gateway

Sits inline on every MCP-based integration. Evaluates each agent request against context, intent, and risk in real time. Can audit, throttle, and terminate sessions before exfiltration completes — no human SOC required.



Runtime Identity + PBAC

Closes the point-in-time trust gap. Every access request re-evaluated continuously against behavioral context, risk signals from CrowdStrike, Zscaler, and Wiz. Parameterized payment tokens limit agent financial scope by merchant and amount.



Zero-Knowledge Biometrics

Keyless acquisition brings server-side, privacy-preserving biometrics trained specifically against synthetic media. Defeats real-time face-swap and voice-cloning attacks. Device-less — works in call centers, factory floors, regulated environments.

THREAT-TO-CAPABILITY MAPPING

Threat Vector	Attack Pathway	Ping Mitigation	Business Risk Closed
Autonomous exploit development	System compromise before patch	Runtime Identity + PBAC	Data breach, ransomware
AI-personalized phishing	Credential theft at scale	PingOne Protect + step-up MFA	Account takeover
Real-time deepfake impersonation	Authorized fraudulent transactions	Keyless ZK Biometrics + Verify	Wire fraud, unauthorized access
Synthetic identity onboarding	Fraud at customer enrollment	PingOne Verify + ZK Biometrics	AML liability, account fraud
Agent credential impersonation	Unattributable agent actions	Agent IAM Core + RFC 8693	Audit failure, misattribution
MCP token hijacking	Lateral movement via agent token	Agent Gateway behavioral enforcement	Data exfiltration
Agent privilege scope creep	Domain-admin equivalent access	PBAC continuous authorization	Lateral movement, abuse
AI-authorized fraudulent payment	Out-of-policy agent transaction	CIBA human-in-the-loop + scoped tokens	Financial fraud, fiduciary risk

IMPLEMENTATION ROADMAP

Four phases to Mythos-ready posture — sequenced for maximum impact within the proliferation window

01

0–60 Days

Exposure Management

Audit all identities for over-privilege. Map your agent inventory. Prioritize agent reachability over CVSS.

Identity Hygiene, Agent IAM
Core Registration

02

30–90 Days

Verified Trust Deployment

Deploy ZK Biometrics for high-risk journeys. Implement PingOne Verify at onboarding. Enable PBAC for privileged access.

Keyless ZK Biometrics,
PingOne Verify, PingOne
Protect

03

60–120 Days

Agentic Readiness

Deploy Agent Gateway on all MCP integrations. Implement RFC 8693 Token Exchange. Configure CIBA for high-risk agentic actions.

Agent Gateway, OAuth RFC
8693, CIBA push approval

04

90–180 Days

Adaptive Orchestration

Deploy DaVinci for all critical auth journeys. Integrate CrowdStrike, Zscaler, and Wiz signals. Build AI-threat response playbooks.

PingOne DaVinci, Identity
Fabric unified control plane



Adaptive Defense Without Code Changes

When a new threat pattern emerges, your response is measured in minutes and not sprint cycles.



No-Code Orchestration

Modify authentication journeys, add step-up challenges, and adjust risk thresholds without engineering involvement.



Hours-Scale Response

When Mythos-class adversaries adapt their approach, defensive playbooks update in minutes and not the next sprint cycle.



Unified Signal Integration

Identity Fabric aggregates risk signals from CrowdStrike, Zscaler, and Wiz for correlated, real-time threat response.

The New Equilibrium of Digital Trust

**Identity is not a component of the security stack.
In the Mythos era, it is the security stack.**

The proliferation window is open. Organizations that build identity-centric resilience now will have controls in place before AI-speed attacks become routine. Those that wait will face those attacks against infrastructure designed for a threat environment that no longer exists.

NEXT STEPS

1

Identity Exposure Analysis

Complimentary agent inventory, over-privilege profiling, and MCP integration risk scoring.

2

Mythos Readiness Assessment

Customized 90-minute session mapping your current posture to the four-phase roadmap.

3

DaVinci Proof of Value

Adaptive orchestration in your environment — no code required.